

Scalable Offline Monitoring

Matúš Harvan

ABB Corporate Research

Joint work with David Basin, Germano Caronni, Sarah Ereth, Felix Klaedtke, and Heiko Mantel.

Motivation



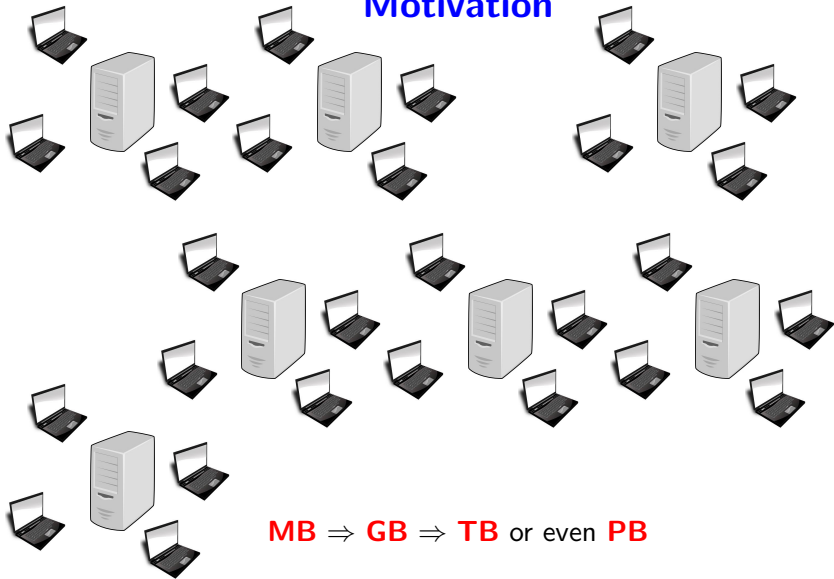
MB

Motivation



MB \Rightarrow GB

Motivation



Contributions

A solution to monitor big data



Contributions



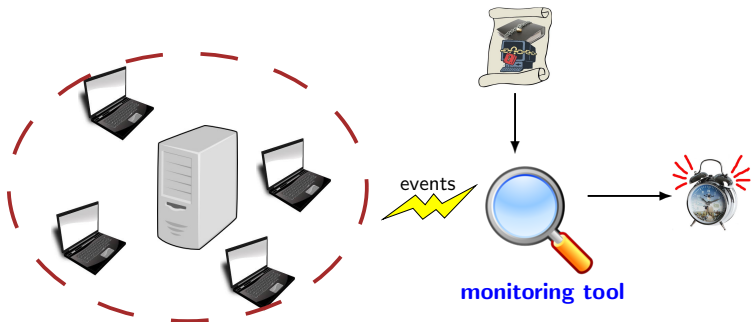
A solution to monitor big data

- ▶ Framework for parallelizing monitoring
 - Operators to slice logs.
 - Soundness & completeness guarantees.
- ▶ Algorithmic realization with MapReduce.
- ▶ Evaluation on real-world data.

Setting and Prior Work



Setting and Prior Work



- ▶ General solution using metric first-order temporal logic (MFOTL).
- ▶ Monitoring algorithm implemented in Monpoly tool.

Example

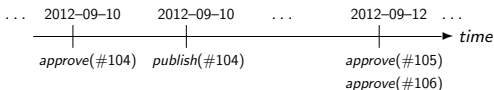


► Policy:

1. Reports must be approved before they are published.
2. Approvals must happen at most 10 days before publication.

► Events are logged with time stamps:

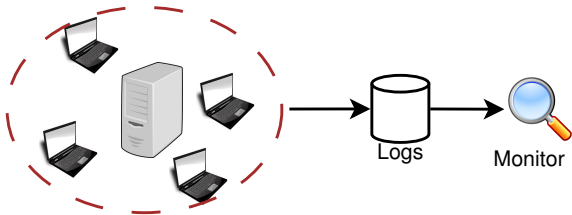
⋮	⋮
2012-09-10	approve (#104)
2012-09-10	publish (#104)
⋮	⋮
2012-09-12	approve (#105)
	approve (#106)
⋮	⋮



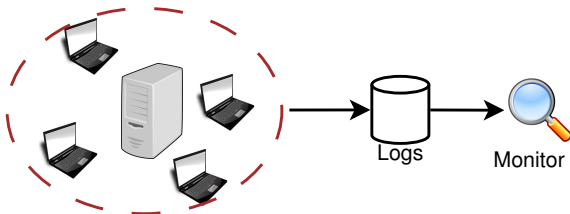
► Policy formalized in metric first-order temporal logic:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

Framework to parallelize monitoring

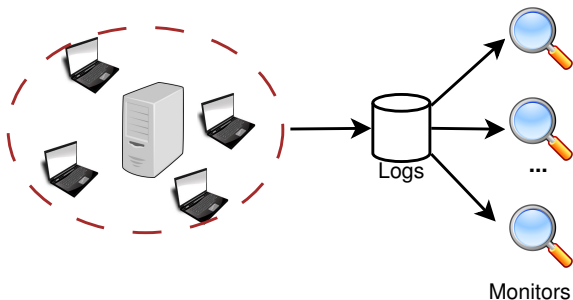


Framework to parallelize monitoring



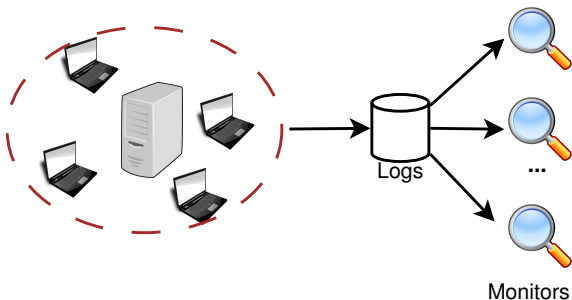
Framework to parallelize monitoring

- Split a large log into smaller log parts ([slices](#)) that can be monitored independently of each other.



Framework to parallelize monitoring

- ▶ Split a large log into smaller log parts (**slices**) that can be monitored independently of each other.
- ▶ Slicing methods:
 1. By time
 2. By data
 3. Filtering
- ▶ Compositionality
- ▶ Sound & complete



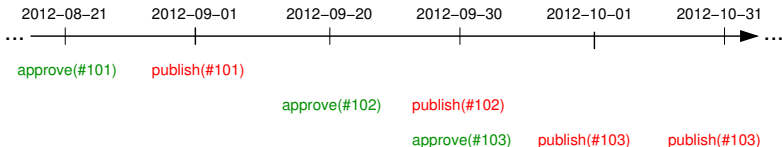
Slicing Time

- Split log based on the timestamps.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different months.



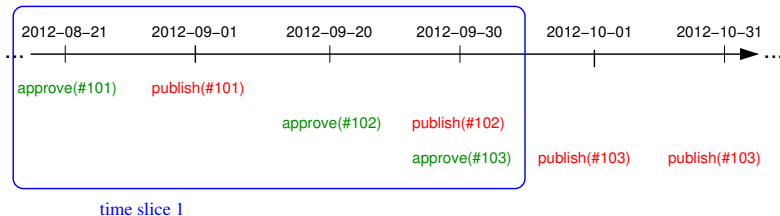
Slicing Time

- Split log based on the timestamps.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different months.



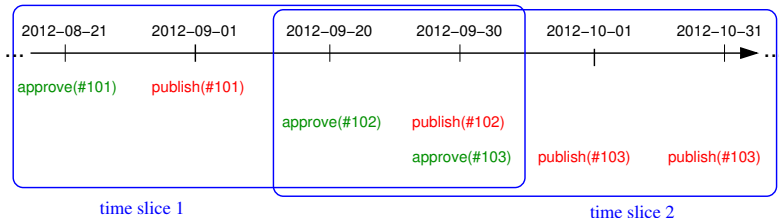
Slicing Time

- Split log based on the timestamps.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different months.



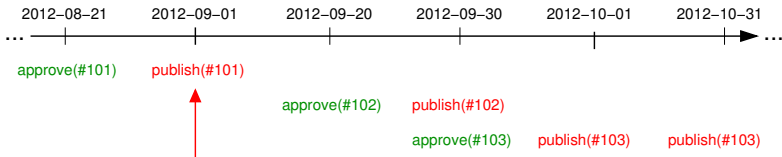
Slicing Time

- Split log based on the timestamps.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different months.



- **Relative interval** determines required additional points.
→ Derived from temporal operators.
- **Restrictions** determine which violations are valid.

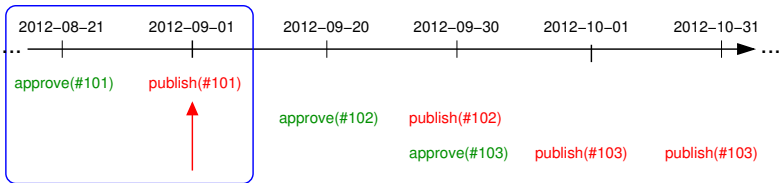
Slicing Time

- Split log based on the timestamps.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different months.



- **Relative interval** determines required additional points.
→ Derived from temporal operators.
- **Restrictions** determine which violations are valid.

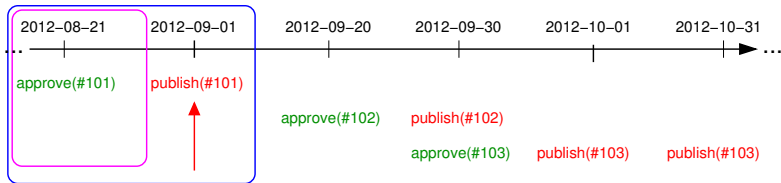
Slicing Time

- Split log based on the timestamps.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different months.



- **Relative interval** determines required additional points.
→ Derived from temporal operators.
- **Restrictions** determine which violations are valid.

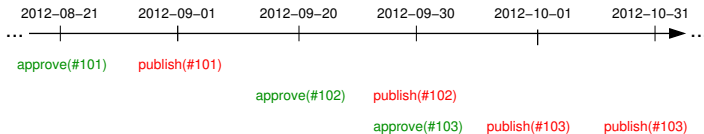
Slicing Data

- Split log based on parameters of log events.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different reports.



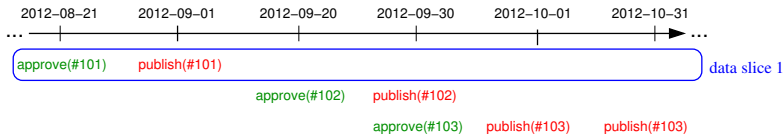
Slicing Data

► Split log based on parameters of log events.

► Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different reports.



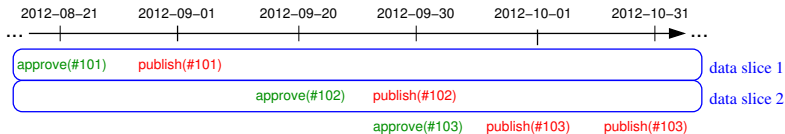
Slicing Data

► Split log based on parameters of log events.

► Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different reports.



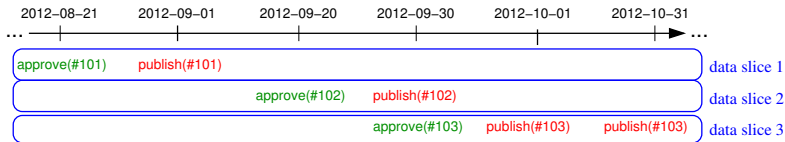
Slicing Data

- Split log based on parameters of log events.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different reports.



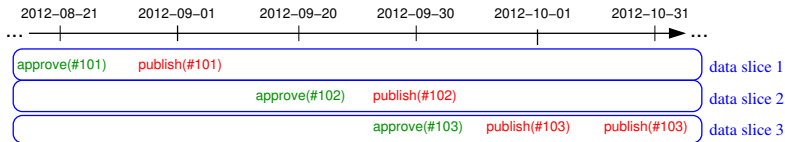
Slicing Data

- Split log based on parameters of log events.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$

- Slices cover different reports.



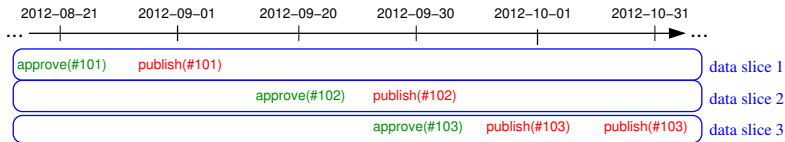
- What about these formulas?

- $\forall r. \text{publish}(r) \rightarrow \blacklozenge \neg \exists r'. \text{publish}(r') \wedge r' > r$
- $\forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 7} \text{publish}(\text{summary})$

Slicing Data

- Split log based on parameters of log events.

- Example: $\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$
 - Slices cover different reports.



- What about these formulas?

- $\forall r. \text{publish}(r) \rightarrow \blacklozenge \neg \exists r'. \text{publish}(r') \wedge r' > r$
- $\forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 7} \text{publish}(\text{summary})$

- How to do data slicing in general?

- Trade-off: efficiency versus size.
- Determine slice membership by inspecting log events independently.

Slicing Data



- ▶ Choose **slicing variable**.
- ▶ Split domain into slicing sets.
→ Each slicing set induces a slice.
- ▶ Criteria whether log event $e(p_1, \dots, p_n)$ belongs into slice:

For each parameter p_i , look for atomic subformula $e(x_1, \dots, x_n)$ where x_i is...

Slicing Data



- Choose **slicing variable**.
- Split domain into slicing sets.
→ Each slicing set induces a slice.
- Criteria whether log event $e(p_1, \dots, p_n)$ belongs into slice:

For each parameter p_i , look for atomic subformula $e(x_1, \dots, x_n)$ where x_i is...

1. the **slicing variable** and p_i is in slicing set,
Example: $\forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$

Slicing Data



- ▶ Choose **slicing variable**.
- ▶ Split domain into slicing sets.
→ Each slicing set induces a slice.
- ▶ Criteria whether log event $e(p_1, \dots, p_n)$ belongs into slice:

For each parameter p_i , look for atomic subformula $e(x_1, \dots, x_n)$ where x_i is...

1. the **slicing variable** and p_i is in slicing set,
Example: $\forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$
2. **another variable**,
Example: $\forall r. \text{publish}(r) \rightarrow \blacklozenge \neg \exists r'. \text{publish}(r') \wedge r' > r$

Slicing Data



- Choose **slicing variable**.
- Split domain into slicing sets.
→ Each slicing set induces a slice.
- Criteria whether log event $e(p_1, \dots, p_n)$ belongs into slice:

For each parameter p_i , look for atomic subformula $e(x_1, \dots, x_n)$ where x_i is...

1. the **slicing variable** and p_i is in slicing set,
Example: $\forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$
2. **another variable**,
Example: $\forall r. \text{publish}(r) \rightarrow \blacklozenge \neg \exists r'. \text{publish}(r') \wedge r' > r$
3. or a **constant**.
Example: $\forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 7} \text{publish}(\text{summary})$

Choice of Slicing Variable Matters

- ▶ Policy: $\Box \forall c. \forall s. \text{login}(c, s) \rightarrow \Diamond_{\leq 6} \text{notify}(0, s)$
- ▶ Log: $\text{login}(1, 1), \text{login}(2, 2), \text{login}(3, 3), \text{login}(4, 4)$
 $\text{notify}(0, 1), \text{notify}(0, 2), \text{notify}(0, 3), \text{notify}(0, 4)$
- ▶ Slicing by variable c
 - Slicing set: $\{1, 2\}$
Slice: $\text{login}(1, 1), \text{login}(2, 2),$
 $\text{notify}(0, 1), \text{notify}(0, 2), \text{notify}(0, 3), \text{notify}(0, 4)$
 - Slicing set: $\{3, 4\}$
Slice: $\text{login}(3, 3), \text{login}(4, 4),$
 $\text{notify}(0, 1), \text{notify}(0, 2), \text{notify}(0, 3), \text{notify}(0, 4)$

Choice of Slicing Variable Matters

- ▶ Policy: $\Box \forall c. \forall s. \text{login}(c, s) \rightarrow \Diamond_{\leq 6} \text{notify}(0, s)$
- ▶ Log: $\text{login}(1, 1), \text{login}(2, 2), \text{login}(3, 3), \text{login}(4, 4)$
 $\text{notify}(0, 1), \text{notify}(0, 2), \text{notify}(0, 3), \text{notify}(0, 4)$
- ▶ Slicing by variable c
 - Slicing set: $\{1, 2\}$
Slice: $\text{login}(1, 1), \text{login}(2, 2),$
 $\text{notify}(0, 1), \text{notify}(0, 2), \text{notify}(0, 3), \text{notify}(0, 4)$
 - Slicing set: $\{3, 4\}$
Slice: $\text{login}(3, 3), \text{login}(4, 4),$
 $\text{notify}(0, 1), \text{notify}(0, 2), \text{notify}(0, 3), \text{notify}(0, 4)$
- ▶ Slicing by variable s
 - Slicing set: $\{1, 2\}$
Slice: $\text{login}(1, 1), \text{login}(1, 2),$
 $\text{notify}(0, 1), \text{notify}(0, 2),$
 - Slicing set: $\{3, 4\}$
Slice: $\text{login}(3, 3), \text{login}(4, 4),$
 $\text{notify}(0, 3), \text{notify}(0, 4)$

Filtering

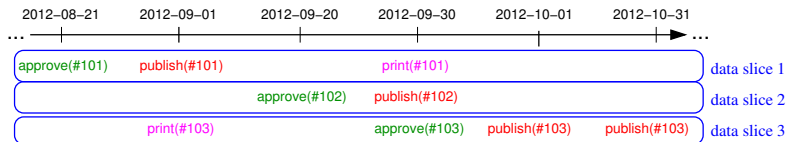
- Discards log parts “irrelevant” to the formula:

1. “Irrelevant” *log events*.

- Already discarded by data slicing.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$



Filtering

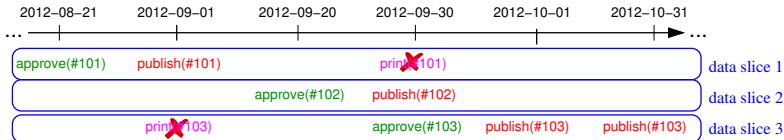
- Discards log parts “irrelevant” to the formula:

1. “Irrelevant” *log events*.

- Already discarded by data slicing.

- Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$



Filtering

► Discards log parts “irrelevant” to the formula:

1. “Irrelevant” *log events*.

- Already discarded by data slicing.

2. *Empty time points*.

- Can safely be done only for some formulas.

Problematic subformula examples:

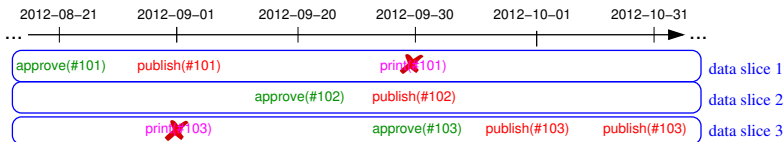
$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 5} \blacklozenge_{\leq 5} \text{approve}(r)$$

$$\Box \forall r. \text{publish}(r) \rightarrow \Box_{\leq 60} \neg \text{product_launch}(r)$$

- Safe formulas approximated by a syntactically-defined fragment.

► Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$



Filtering

► Discards log parts “irrelevant” to the formula:

1. “Irrelevant” *log events*.

- Already discarded by data slicing.

2. *Empty time points*.

- Can safely be done only for some formulas.

Problematic subformula examples:

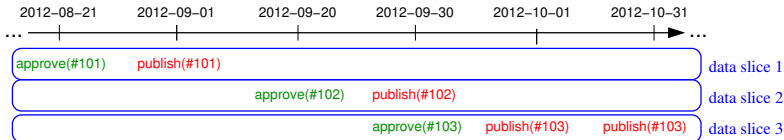
$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 5} \blacklozenge_{\leq 5} \text{approve}(r)$$

$$\Box \forall r. \text{publish}(r) \rightarrow \Box_{\leq 60} \neg \text{product_launch}(r)$$

- Safe formulas approximated by a syntactically-defined fragment.

► Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$



Filtering

► Discards log parts “irrelevant” to the formula:

1. “Irrelevant” *log events*.

- Already discarded by data slicing.

2. *Empty time points*.

- Can safely be done only for some formulas.

Problematic subformula examples:

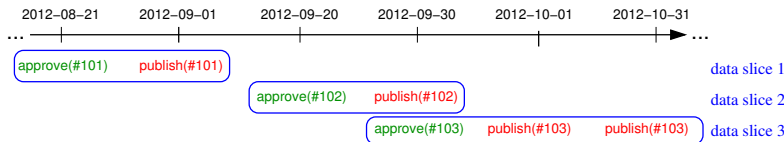
$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 5} \blacklozenge_{\leq 5} \text{approve}(r)$$

$$\Box \forall r. \text{publish}(r) \rightarrow \Box_{\leq 60} \neg \text{product_launch}(r)$$

- Safe formulas approximated by a syntactically-defined fragment.

► Example:

$$\Box \forall r. \text{publish}(r) \rightarrow \blacklozenge_{\leq 10} \text{approve}(r)$$



Composition

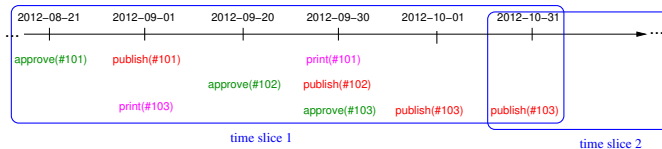
- Combine slicing methods arbitrarily.



- Sound & complete by construction.

Composition

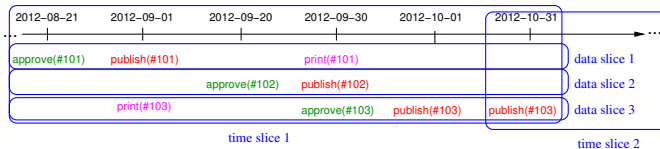
- Combine slicing methods arbitrarily.
 - Slice time



- Sound & complete by construction.

Composition

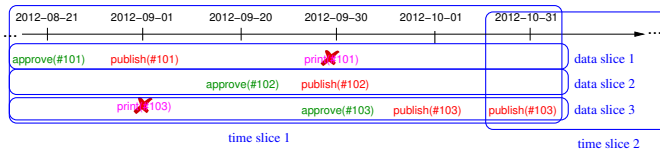
- Combine slicing methods arbitrarily.
- Slice time
- Slice data



- Sound & complete by construction.

Composition

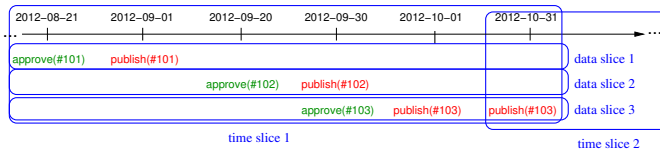
- Combine slicing methods arbitrarily.
 - Slice time
 - Slice data
 - Filter “irrelevant” log events.



- Sound & complete by construction.

Composition

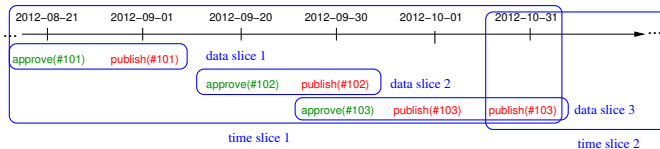
- Combine slicing methods arbitrarily.
 - Slice time
 - Slice data
 - Filter “irrelevant” log events.



- Sound & complete by construction.

Composition

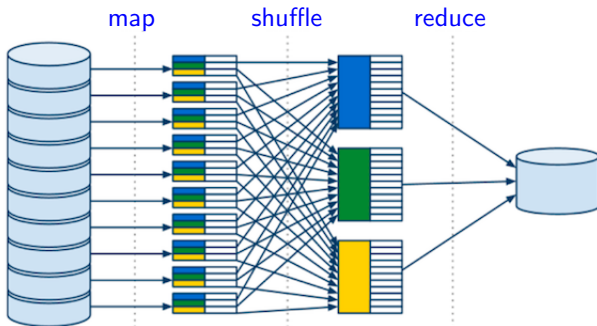
- ▶ Combine slicing methods arbitrarily.
- Slice time
- Slice data
- Filter “irrelevant” log events.
- Filter empty time points.



- ▶ Sound & complete by construction.

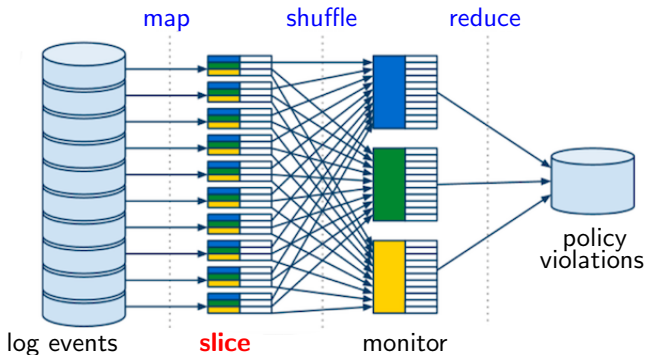
Algorithmic Realization with MapReduce

- MapReduce framework supports data-intensive distributed computations



Algorithmic Realization with MapReduce

- ▶ MapReduce framework supports data-intensive distributed computations

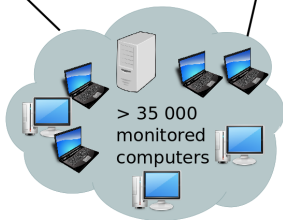


- ▶ Advantages:
 - Computation distribution
 - Load-balancing
 - Fault-tolerance

Google Case Study

sensitive resources

update server



syslog
messages



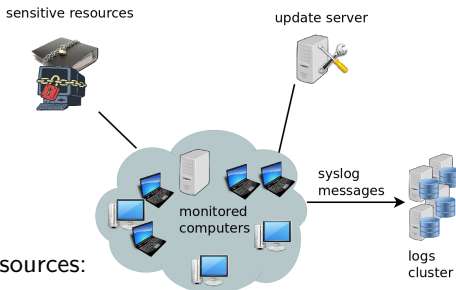
logs
cluster

► Verify hypotheses:

1. Approach scales to very large logs.
2. Slicers suitable for real-world policies.

- Logged 1 TB per day.
- Extracted log events:
 - Cover approximately 2 years.
 - 26 billion log events
 - 0.4 TB

Policies



1. Authentication to access sensitive resources:

P1: Entering credentials takes at least 1 second.

P2: Only from updated computers (updated within the last 3 days).

2. SSH sessions:

P3: Not longer than 24 hours.

3. System updates:

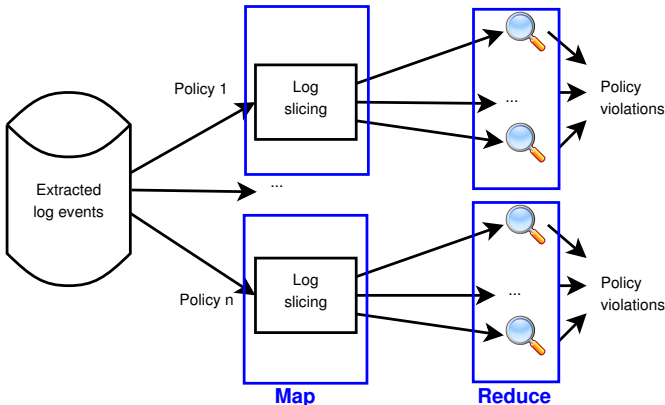
P4: Update at least every 3 days.

P5: After downloading updates, apply them within 30 minutes.

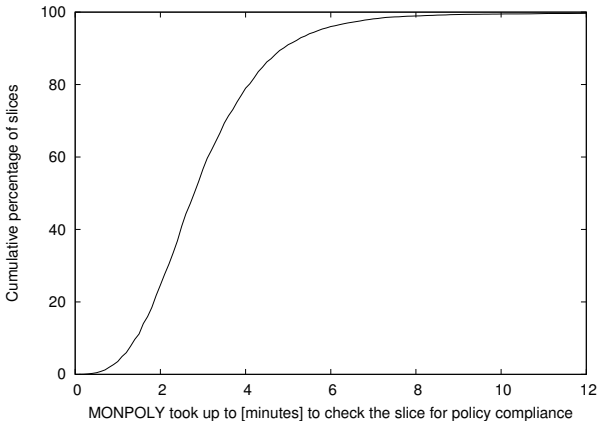
P6: If a computer claims to be up-to-date, it must have updated during the last 24 hours.

Evaluation

- ▶ Log split into 10,000 data slices (by computer logging the event).
- ▶ For each policy, slices monitored by 1000 computers in parallel.
- ▶ Monitoring finished in 2.5 hrs, except for policy *P3* (12 hrs).

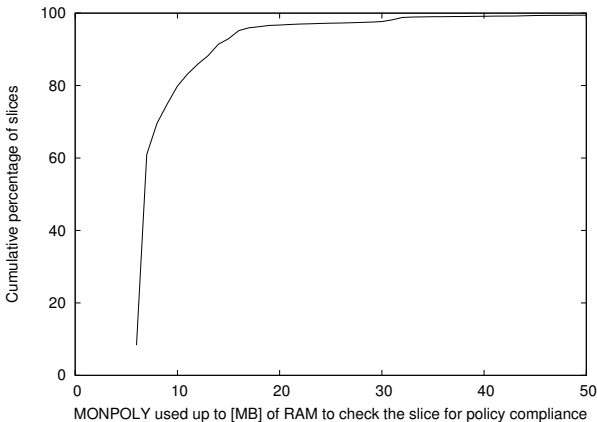


Time to Check a Single Slice



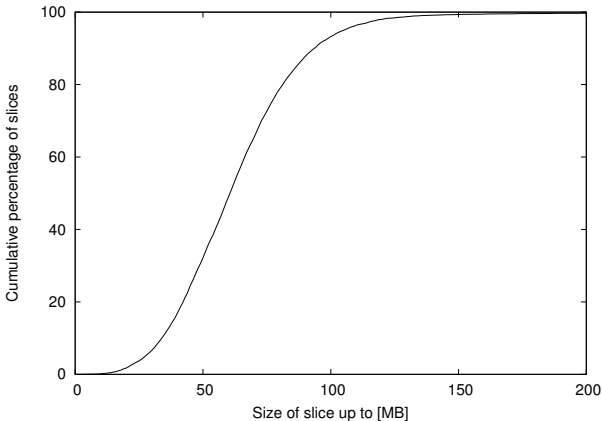
- ▶ 90% of slices monitored within 5 min.
- ▶ Median time: 3 min.
- ▶ Max time: 11 hours (policy *P3*).

Memory to Check a Single Slice



- ▶ 90% of slices monitored with less than 14 MB of memory.
- ▶ Median memory: 6–10 MB.
- ▶ Max memory: 510 MB (policy *P3*).

Distribution of Slice Sizes



- ▶ 90% of slices smaller than 94 MB.
- ▶ Median size of a slice: 61 MB.
- ▶ Three slices over 1 GB, max 1.8 GB.

Related Work

- ▶ B. Barre, M. Klein, M. Soucy-Boivin, P.-A. Ollivier, and S. Hallé.
[MapReduce for parallel trace validation of LTL properties.](#)
RV 2012.
 - Parallelization based on formula structure rather than data.
→ Limited scalability
 - Case study: < 5 million log events, monitored on a single computer.

- ▶ G. Rosu and F. Chen.
[Semantics and algorithms for parametric monitoring.](#)
Log. Method. Comput. Sci., 2012.
 - Parametric temporal logic less expressive than first-order temporal logic.
 - Case study: 155 million log events, monitored on a single computer.
 - No distributed computation.

Conclusions

► Scalable monitoring solution

- Framework to parallelize monitoring.
- Algorithmic realization.
- Case study.

► Very large case study (largest to date?) of applying runtime verification techniques to big data.

► Future work

- Overcome offline limitation: use a stream processing platform.
- Overcome unbalanced distribution of log slices.
- Design and evaluate other slicing operators.



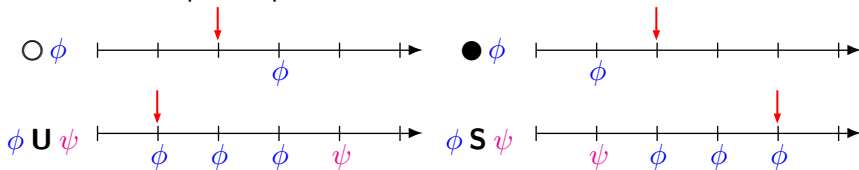
Backup slides

Metric First-Order Temporal Logic

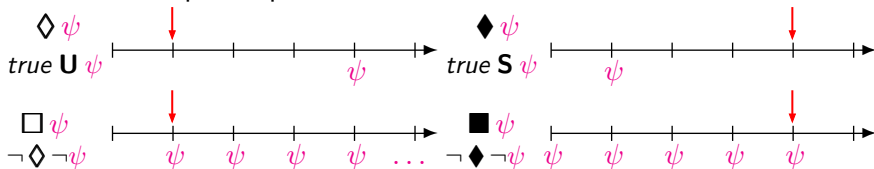
- ▶ **First-order** for expressing relations on data
 - $<, =$
 - \neg, \wedge, \vee
 - \forall, \exists
- ▶ Metric **temporal operators** for expressing qualitative and quantitative timing information
 - Past: $\bullet, \mathbf{S}, \blacklozenge, \blacksquare$
(PREVIOUS, SINCE, ONCE, ALWAYS PAST)
 - Future: $\circ, \mathbf{U}, \lozenge, \square$
(NEXT, UNTIL, EVENTUALLY, ALWAYS)
 - Metric operators add timing constraints,
e.g., $\blacklozenge_{\leq 10}$ for once within 10 time units
(ONCE $[0, 11)$)

Standard linear temporal operators

► Primitive temporal operators

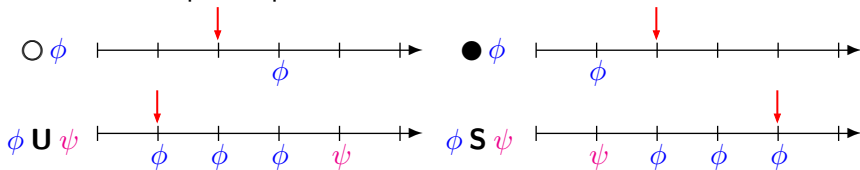


► Derived temporal operators

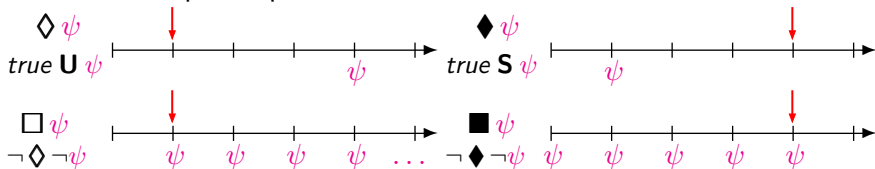


Metric temporal operators

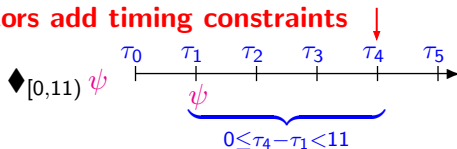
Primitive temporal operators



Derived temporal operators



Metric operators add timing constraints



Monpoly

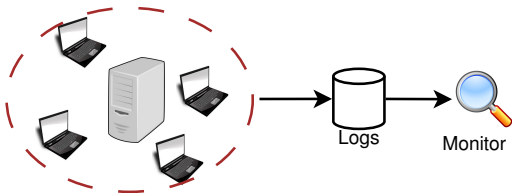


- ▶ Monitoring tool to determine which events in a log violate given rules.
 - Input: policy, log
 - Output: violations
- ▶ Policies specified in metric first-order temporal logic.
- ▶ Result of multi-year project at ETH Zurich.

Monpoly

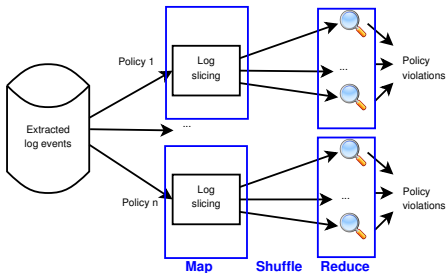


- ▶ Monitoring tool to determine which events in a log violate given rules.
 - Input: policy, log
 - Output: violations
- ▶ Policies specified in metric first-order temporal logic.
- ▶ Result of multi-year project at ETH Zurich.
- ▶ Single-threaded, not parallelized.



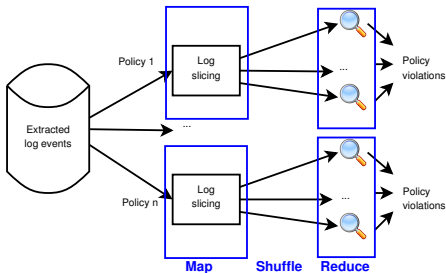
MapReduce Details

- **Map:** Split log into slices, emitting each log event with:
 - Primary key: log slice identifier
 - Secondary key: timestamp



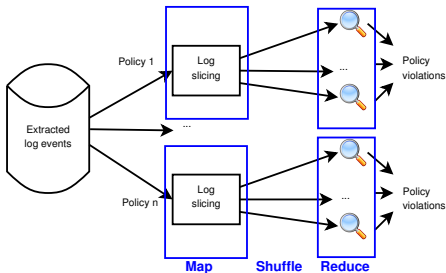
MapReduce Details

- ▶ **Map:** Split log into slices, emitting each log event with:
 - Primary key: log slice identifier
 - Secondary key: timestamp
- ▶ **Shuffle:** Collect all parts of a log slice.



MapReduce Details

- ▶ **Map:** Split log into slices, emitting each log event with:
 - Primary key: log slice identifier
 - Secondary key: timestamp
- ▶ **Shuffle:** Collect all parts of a log slice.
- ▶ **Reduce:** Monitor a log slice.
 - Monpoly runs in a child process.
 - Log slice piped into Monpoly.
 - Monpoly's output (policy violations) returned by reducer.



Logs

- ▶ Raw logs: 1TB of syslog messages per day
- ▶ Extracted log events:
 - Cover approximately 2 years.
 - 26 billion log events
 - 0.4 TB

